

AI Village: real evidence and a solved reference-recovery task

Observational evidence

The exact source is dataset-1ac43f5141de, version 1, fingerprint 16849167aecaa4815684f621d12940c905f18fa327e615778dae68362ee4883e. It retains 2,000 posts, September 8-11, 2025: 1,983 agent and 17 user posts. The requested September 8-22 period is right-censored by that limit. This is neither exhaustive discovery nor a prevalence estimate.

The messages report failed URLs, access tests and reference changes. Literal source search now retrieves actual matching posts with original offsets; it cannot authenticate historical Google state or infer readers. A real investigator's candidate behavior-328c91628012 v1 was skeptically rejected as a distinct mechanism; fingerprint 356d54274e224cb86b4f406dbcc0f870f622c36029b5af11a3ec8cdb8e31dcc5. Ordinary URL, session and shared-artifact explanations remain unresolved. Novelty is not established.

New controlled contract

The family is `single_document_reference_repair`. Two fresh roles act on one original IRB document. Content, viewer ACL and initial authenticated profiles are held fixed. The auditor starts with one broken reference and must actually open the current original. Both conditions have this same explicit goal and the same ten legal tools. Title search, exact navigation and direct peer messages are available; recreation and sharing are excluded. The pre-treatment navigation probe is code-generated experimental setup, not a model action or historical event.

The assigned notes are neutral versus canonical-reference checking. Two matched pairs use the same generated starting world, isolated fresh teams and randomized within-pair execution order. The primary integer `verified_repaired_reference` requires executed auditor opening plus correct current content and authorized profile. Separate inspection or message-receipt gates are not added to that outcome. The proxy omits historical Google services/state and original policies; historical equivalence is unestablished.

Completed real run

Recorded measure	Neutral note	Canonical-reference check
Fresh two-role teams	2	2
Verified repaired references	2 / 2	2 / 2
Real subject decisions	32	32
Recorded reference failures	0	0

All four teams solved the registered task within the 16-decision budget. Their first successful auditor opens occurred at recorded steps 11, 2, 2 and 1 respectively. Every owner queued an actual canonical reference; this does not prove the teammate read it or identify the route causing success. Only one of four auditors produced the separate current-content inspection receipt. Opening, inspection and comprehension remain distinct.

The paired primary difference is **0 percentage points**, with saved 95% interval **-100 to +100 percentage points**, method `bounded_pair_Hoeffding_95`. Unit: matched two-role-team difference, two pairs. The exact paired sign-flip test has $p = 1$ with four sign assignments under its sharp null. Both arms succeeding demonstrates task feasibility and executable access recovery, not treatment benefit, equivalence or a historical mechanism. The earlier broader four-document task's zero result remains a separate negative task/goal diagnostic; it is not pooled or rescored.

Exact record, all version 1	Payload fingerprint
village_recovery_experiment-63d316aa7898	c0a78378267faff180b1ad6306c0752d9df3595357b267de9c880b03ad51cc38
guided_result-20ccadcf545c	beec60f72420bb2434c083c2e17e0672eefbe8c1e3146c86c8b2a437d5624fa7
guided_plan-9f11f66e8b9b	4abb454039219d57e19d624cd072a26f8826681975de713fed1c0800b8a8f4de
guided_simulator-recovery-4abb454039219d57	27694d8c3729dc9cda9474b1c521f8dcb374f4e01999f7f425949415c4a74d60
verification-43d7bf1ace43	8a4d2abdcce1d9372ff76165fa7e129d63fc700a1e187e39aa270c774c77cfac
claim_audit-5362cc6f27dc	22d42967a06fa2c544a6867bdd5b7a65d4ac65df3b7ab4dc3b53773efc5e7714


Fresh independent replay passed. All 64 unique actual provider responses matched retained payload hashes, model actions and world transitions. Archive, source, job, proof and finite fact records agree. These are local integrity/execution checks, not signed provider attestation, semantic validation or evidence of historical equivalence. The deterministic fact checker is not an AI evaluator's causal interpretation.

The native figure uses only this exact result: paired binary outcomes, saved interval, descriptive reference/failure counts and ordinal tool progression. Unknown measurements remain unknown; action positions are not durations. No studies are pooled.

Local Laya inference was blocked by the memory guard. Optional Laya/Jev adapters are not evidence of classifier accuracy or calibration in this investigation.

Project story · Demonstration.

Executed single-document reference recovery




Four fresh two-role teams, two matched pairs, 64 real subject decisions. Both notes solved 2/2 verified repaired references. The saved paired 95% interval spans -100 to +100 percentage points. Secondary counts and ordinal action histories are descriptive; they do not identify a behavioral mechanism, historical cause or demonstrated note benefit.

Inside the real AI Village logs


1,983 agent posts · 17 user posts excluded · Sep 8–11, 2025

AI Village: when the agents posted



1,983 recorded agent posts across six agents, September 8–11. Empty time bins describe retained chat, not agent inactivity.

Agent names appearing together in recorded posts



Edges = agent-authored posts naming both agents · only five largest edge counts labeled
All edge witnesses remain in the JSON. Audience, receipt and influence are unknown.

Lines count posts naming both agents. GPT-5 and o3 appear together in 113 posts. This lexical graph does not establish message delivery, reading or influence.